

Computer Vision

Lecture Set 04 Image Features

Huei-Yung Lin

lin@ee.ccu.edu.tw

Robot Vision Lab
Department of Computer Science and Information Engineering
National Taipei University of Technology
Taipei 106, Taiwan

October 23, 2022

Announcements - 10/24/22

- Homework 2 will be given later, due in two weeks (11/7).
- Exam I is scheduled on 11/7, covers up to Lecture Set 04.

Image Features?



Image Features?




Image Features

- Image features in computer vision
 - ▶ **Global feature** – A global property of an image. e.g. the average grey value, the area in pixel, etc.
 - ▶ **Local feature** – A part of the image with some special properties. e.g. a circle, a line, a textured regions, etc.
- Image features (a practical definition): *local meaningful detectable parts of the image*
 - ▶ Points
 - ▶ Edges: step edges, line edges
 - ▶ Contours: closed contours are boundaries
 - ▶ Regions: similar color, similar feature, etc.

Image Features

- Image features in computer vision
 - ▶ **Global feature** – A global property of an image. e.g. the average grey value, the area in pixel, etc.
 - ▶ **Local feature** – A part of the image with some special properties. e.g. a circle, a line, a textured regions, etc.
- Image features (a practical definition): *local meaningful detectable parts of the image*
 - ▶ Points
 - ▶ Edges: step edges, line edges
 - ▶ Contours: closed contours are boundaries
 - ▶ Regions: similar color, similar feature, etc.

Remarks

- Most vision systems begin by detecting and locating some features in the input images
- In 3-D vision, feature extraction is an intermediate step, not the goal
 - ▶ We do not extract lines just to obtain line maps
 - ▶ We extract lines to navigate robots in corridors or for camera calibration
- Feature extraction is for certain purpose of the system
- “Perfect” feature extraction is not necessary (and also not possible?)

Remarks

- Most vision systems begin by detecting and locating some features in the input images
- In 3-D vision, feature extraction is an intermediate step, not the goal
 - ▶ We do not extract lines just to obtain line maps
 - ▶ We extract lines to navigate robots in corridors or for camera calibration
- Feature extraction is for certain purpose of the system
- “Perfect” feature extraction is not necessary (and also not possible?)

Remarks

- Most vision systems begin by detecting and locating some features in the input images
- In 3-D vision, feature extraction is an intermediate step, not the goal
 - ▶ We do not extract lines just to obtain line maps
 - ▶ We extract lines to navigate robots in corridors or for camera calibration
- Feature extraction is for certain purpose of the system
- “Perfect” feature extraction is not necessary (and also not possible?)

Edge Detection

- Edges typically occur on the boundary between two different regions in an image
- Edge detection is frequently the first step in recovering information from images
- An edge is a significant local change in the intensity
 - ▶ Usually associated with a discontinuity in either the image intensity or its first derivative
 - ▶ The discontinuities can be step or line (ideally)
 - ▶ In reality, step becomes ramp and line becomes roof due to the smoothing of sharp edges in most images

Edge Detection

- Edges typically occur on the boundary between two different regions in an image
- Edge detection is frequently the first step in recovering information from images
- An edge is a significant local change in the intensity
 - ▶ Usually associated with a discontinuity in either the image intensity or its first derivative
 - ▶ The discontinuities can be step or line (ideally)
 - ▶ In reality, step becomes ramp and line becomes roof due to the smoothing of sharp edges in most images



Step



Ramp




Line



Roof


Edge Detection Scheme

- Edges happen at places where the image values exhibit sharp variation



Edge Detection Scheme

- The basic approach to edge detection is to compute “spatial derivatives” of the intensity image
- The act of taking spatial derivatives is usually approximated by convolution



Edge = sharp variation



Large first derivative

Gradient

- The gradient is a measure of change in a function
- Significant changes in gray values can be detected by using a “discrete approximation” to the gradient
- The gradient is a 2-D equivalent of the first derivative and is defined as a vector

$$\mathbf{G}[f(x, y)] = \begin{bmatrix} G_x \\ G_y \end{bmatrix} = \begin{bmatrix} \partial f / \partial x \\ \partial f / \partial y \end{bmatrix}$$

- ▶ The vector $\mathbf{G}[f(x, y)]$ points in the direction of the maximum rate of increase of the function $f(x, y)$
- ▶ The magnitude of the gradient given by $|\mathbf{G}[f(x, y)]| = \sqrt{G_x^2 + G_y^2}$ equals the maximum rate of increase of $f(x, y)$ per unit distance

Gradient

- The gradient is a measure of change in a function
- Significant changes in gray values can be detected by using a “discrete approximation” to the gradient
- The gradient is a 2-D equivalent of the first derivative and is defined as a vector

$$\mathbf{G}[f(x, y)] = \begin{bmatrix} G_x \\ G_y \end{bmatrix} = \begin{bmatrix} \partial f / \partial x \\ \partial f / \partial y \end{bmatrix}$$

- ▶ The vector $\mathbf{G}[f(x, y)]$ points in the direction of the maximum rate of increase of the function $f(x, y)$
- ▶ The magnitude of the gradient given by $|\mathbf{G}[f(x, y)]| = \sqrt{G_x^2 + G_y^2}$ equals the maximum rate of increase of $f(x, y)$ per unit distance

Approximation of Gradient

- The absolute values are commonly used to approximate the gradient magnitude

$$|\mathbf{G}[f(x, y)]| \approx |G_x| + |G_y| \quad \text{or} \quad |\mathbf{G}[f(x, y)]| \approx \max(|G_x|, |G_y|)$$

- The direction of gradient is define as

$$\alpha(x, y) = \tan^{-1} \left(\frac{G_x}{G_y} \right)$$

- The magnitude of the gradient is “independent” of the direction of the edge
- Such operators are called **isotropic operators**

Digital Approximation

- For digital images, the gradient approximation can be

$$G_x \approx f[i, j + 1] - f[i, j] \quad G_y \approx f[i, j] - f[i + 1, j]$$

Digital Approximation

- Simple convolution masks:

- ▶ 2×1 and 1×2 :

| | |
|----|---|
| -1 | 1 |
|----|---|

| |
|----|
| -1 |
| 1 |

- ▶ 2×2 and 2×2 :

| | |
|----|---|
| -1 | 1 |
| -1 | 1 |

| | |
|----|----|
| 1 | 1 |
| -1 | -1 |

- ▶ 3×1

$$\frac{df(x)}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$
$$\frac{df(x)}{dx} \approx \frac{f(x + 1) - f(x - 1)}{2}$$

Convolve with:

| | | |
|----|---|---|
| -1 | 0 | 1 |
|----|---|---|

Edge Detection

- 1D Case

- $I(x)$

- $\frac{dI(x)}{dx}$

- $\left| \frac{dI(x)}{dx} \right| > \text{threshold}$

- No orientation

- 2D

- $I(x, y)$

- $\nabla I(x, y) = \begin{pmatrix} \frac{\partial I(x, y)}{\partial x} \\ \frac{\partial I(x, y)}{\partial y} \end{pmatrix} = \begin{pmatrix} I_x(x, y) \\ I_y(x, y) \end{pmatrix}$

- $|\nabla I(x, y)| = \sqrt{I_x(x, y)^2 + I_y(x, y)^2} > \text{threshold}$

- $\tan \theta = \frac{I_x(x, y)}{I_y(x, y)}$

Edge Detection


- 1D Case

- $I(x)$

- $\frac{dI(x)}{dx}$

- $\left| \frac{dI(x)}{dx} \right| > \text{threshold}$

- No orientation



- 2D

- $I(x, y)$

- $\nabla I(x, y) = \begin{pmatrix} \frac{\partial I(x, y)}{\partial x} \\ \frac{\partial I(x, y)}{\partial y} \end{pmatrix} = \begin{pmatrix} I_x(x, y) \\ I_y(x, y) \end{pmatrix}$

- $|\nabla I(x, y)| = \sqrt{I_x(x, y)^2 + I_y(x, y)^2} > \text{threshold}$

- $\tan \theta = \frac{I_x(x, y)}{I_y(x, y)}$

Edge Detection


- 1D Case

- $I(x)$

- $\frac{dI(x)}{dx}$

- $\left| \frac{dI(x)}{dx} \right| > \text{threshold}$

- No orientation




- 2D

- $I(x, y)$

- $\nabla I(x, y) = \begin{pmatrix} \frac{\partial I(x, y)}{\partial x} \\ \frac{\partial I(x, y)}{\partial y} \end{pmatrix} = \begin{pmatrix} I_x(x, y) \\ I_y(x, y) \end{pmatrix}$

- $|\nabla I(x, y)| = \sqrt{I_x(x, y)^2 + I_y(x, y)^2} > \text{threshold}$

- $\tan \theta = \frac{I_x(x, y)}{I_y(x, y)}$



Edge Detection

- Vertical edges:
 - ▶ Convolve with

| | | |
|----|---|---|
| -1 | 0 | 1 |
|----|---|---|

- Horizontal edges:
 - ▶ Convolve with

| |
|----|
| -1 |
| 0 |
| 1 |

The Essential Edge Descriptor

- Edge position or center
 - ▶ The image position at which the edge is located
 - ▶ Usually saved in a binary image (1 : edge, 0 : no edge)
- Edge normal
 - ▶ The direction (unit vector) of the maximum intensity variation at the edge point
 - ▶ This identifies the direction perpendicular to the edge
- Edge direction
 - ▶ The direction perpendicular to the edge normal
 - ▶ This identifies the direction tangent to the edge
- Edge strength
 - ▶ A measure of the local image contrast; i.e., how marked the intensity variation is across the edge along the normal

The Essential Edge Descriptor

- Edge position or center
 - ▶ The image position at which the edge is located
 - ▶ Usually saved in a binary image (1 : edge, 0 : no edge)
- Edge normal
 - ▶ The direction (unit vector) of the maximum intensity variation at the edge point
 - ▶ This identifies the direction perpendicular to the edge
- Edge direction
 - ▶ The direction perpendicular to the edge normal
 - ▶ This identifies the direction tangent to the edge
- Edge strength
 - ▶ A measure of the local image contrast; i.e., how marked the intensity variation is across the edge along the normal

The Essential Edge Descriptor

- Edge position or center
 - ▶ The image position at which the edge is located
 - ▶ Usually saved in a binary image (1 : edge, 0 : no edge)
- Edge normal
 - ▶ The direction (unit vector) of the maximum intensity variation at the edge point
 - ▶ This identifies the direction perpendicular to the edge
- Edge direction
 - ▶ The direction perpendicular to the edge normal
 - ▶ This identifies the direction tangent to the edge
- Edge strength
 - ▶ A measure of the local image contrast; i.e., how marked the intensity variation is across the edge along the normal

The Essential Edge Descriptor

- Edge position or center
 - ▶ The image position at which the edge is located
 - ▶ Usually saved in a binary image (1 : edge, 0 : no edge)
- Edge normal
 - ▶ The direction (unit vector) of the maximum intensity variation at the edge point
 - ▶ This identifies the direction perpendicular to the edge
- Edge direction
 - ▶ The direction perpendicular to the edge normal
 - ▶ This identifies the direction tangent to the edge
- Edge strength
 - ▶ A measure of the local image contrast; i.e., how marked the intensity variation is across the edge along the normal

Edge Detection Steps

- Noise smoothing (filtering)
 - ▶ Suppress noise without destroying the true edges
- Edge enhancement
 - ▶ Design a filter responding to edges
 - ▶ Usually performed by computing gradient magnitude
- Edge localization
 - ▶ Thinning (non-maximum suppression)
 - ▶ Thresholding (used to decide whether the output is an edge point or not)
- Some Assumptions
 - ▶ The edge enhancement filter is linear
 - ▶ The filter must be optimal for noisy step edge
 - ▶ The image noise is additive, white and Gaussian

Edge Detection Steps

- Noise smoothing (filtering)

- ▶ Suppress noise without destroying the true edges

- Edge enhancement

- ▶ Design a filter responding to edges
- ▶ Usually performed by computing gradient magnitude

- Edge localization

- ▶ Thinning (non-maximum suppression)
- ▶ Thresholding (used to decide whether the output is an edge point or not)

- Some Assumptions

- ▶ The edge enhancement filter is linear
- ▶ The filter must be optimal for noisy step edge
- ▶ The image noise is additive, white and Gaussian

Edge Detection Steps

- Noise smoothing (filtering)
 - ▶ Suppress noise without destroying the true edges
- Edge enhancement
 - ▶ Design a filter responding to edges
 - ▶ Usually performed by computing gradient magnitude
- Edge localization
 - ▶ Thinning (non-maximum suppression)
 - ▶ Thresholding (used to decide whether the output is an edge point or not)
- Some Assumptions
 - ▶ The edge enhancement filter is linear
 - ▶ The filter must be optimal for noisy step edge
 - ▶ The image noise is additive, white and Gaussian

Edge Detection Steps


- Noise smoothing (filtering)
 - ▶ Suppress noise without destroying the true edges
- Edge enhancement
 - ▶ Design a filter responding to edges
 - ▶ Usually performed by computing gradient magnitude
- Edge localization
 - ▶ Thinning (non-maximum suppression)
 - ▶ Thresholding (used to decide whether the output is an edge point or not)
- Some Assumptions
 - ▶ The edge enhancement filter is linear
 - ▶ The filter must be optimal for noisy step edge
 - ▶ The image noise is additive, white and Gaussian

Edge Detection Steps

- Noise smoothing (filtering)
 - ▶ Suppress noise without destroying the true edges
- Edge enhancement
 - ▶ Design a filter responding to edges
 - ▶ Usually performed by computing gradient magnitude
- Edge localization
 - ▶ Thinning (non-maximum suppression)
 - ▶ Thresholding (used to decide whether the output is an edge point or not)
- Some Assumptions
 - ▶ The edge enhancement filter is linear
 - ▶ The filter must be optimal for noisy step edge
 - ▶ The image noise is additive, white and Gaussian

Noise Suppression

- The differential kernels act as “high pass filters” which tend to amplify noise
- This is why edge detection is usually preceded by a noise reduction or filtering operation



Noise Smoothing & Edge Detection

- Prewitt Edge Detector (vertical)

- ▶ Convolve with:

| | | |
|----|---|---|
| -1 | 0 | 1 |
| -1 | 0 | 1 |
| -1 | 0 | 1 |

Noise Smoothing

Vertical Edge Detection

- Prewitt Edge Detector (horizontal)

- ▶ Convolve with:

| | | |
|----|----|----|
| -1 | -1 | -1 |
| 0 | 0 | 0 |
| 1 | 1 | 1 |

Horizontal Edge Detection

Noise Smoothing

Roberts Detector

- $G_x :$

| | |
|----|----|
| +1 | 0 |
| 0 | -1 |


- $G_y :$

| | |
|----|----|
| 0 | +1 |
| -1 | 0 |

- $|G| = \sqrt{G_x^2 + G_y^2}$

Roberts Detector

- The Roberts detector gives an approximation to the continuous gradient at $[i + 1/2, j + 1/2]$. (not at $[i, j]$, why?)



Sobel Detector

- Sobel detector gives more weight to the 4-neighbors
- Emphasize the pixels closer to the center of the mask (compare with Prewitt!)

• $G_x :$

| | | |
|----|---|---|
| -1 | 0 | 1 |
| -2 | 0 | 2 |
| -1 | 0 | 1 |


$G_y :$

| | | |
|----|----|----|
| -1 | -2 | -1 |
| 0 | 0 | 0 |
| 1 | 2 | 1 |

• $\theta = \tan^{-1} \frac{G_y}{G_x}$

Sobel Detector

- Sobel operator is one of the most commonly used edge detector
(See handout for how it is derived)



Examples



Original



Sobel



Roberts











Canny



Prewitt

Examples



The Canny Principle

- Canny derived the form of the “optimal” linear filter for detecting edges in 1-D
 - ▶ Noise smoothing
 - ▶ Edge enhancement
 - ▶ Edge localization
- The edge was modeled as a simple step corrupted by additive Gaussian noise
- Experiments consistently show that it performs very well
- Probably the most used by computer vision practitioners

The Canny Principle

- Canny derived the form of the “optimal” linear filter for detecting edges in 1-D
 - ▶ Noise smoothing
 - ▶ Edge enhancement
 - ▶ Edge localization
- The edge was modeled as a simple step corrupted by additive Gaussian noise
- Experiments consistently show that it performs very well
- Probably the most used by computer vision practitioners

Canny's Criteria

- Detection – the important edges should not be missed and there should be no spurious responses
- Localization – the distance between the actual and located position of the edge should be minimal
- Single response – minimize multiple responses to a single edge
 - ▶ This is partly covered by the first criterion (since if there are two responses to a single edge, one of them should be considered as false)
 - ▶ This criterion solves the problem of an edge corrupted by noise
- Canny Edge Detector
 - ▶ Uses a mathematical model of the edge and the noise
 - ▶ Formalizes a performance criteria
 - ▶ Synthesizes the best filter

Canny's Criteria

- Detection – the important edges should not be missed and there should be no spurious responses
- Localization – the distance between the actual and located position of the edge should be minimal
- Single response – minimize multiple responses to a single edge
 - ▶ This is partly covered by the first criterion (since if there are two responses to a single edge, one of them should be considered as false)
 - ▶ This criterion solves the problem of an edge corrupted by noise
- Canny Edge Detector
 - ▶ Uses a mathematical model of the edge and the noise
 - ▶ Formalizes a performance criteria
 - ▶ Synthesizes the best filter

Canny's Criteria

- Detection – the important edges should not be missed and there should be no spurious responses
- Localization – the distance between the actual and located position of the edge should be minimal
- Single response – minimize multiple responses to a single edge
 - ▶ This is partly covered by the first criterion (since if there are two responses to a single edge, one of them should be considered as false)
 - ▶ This criterion solves the problem of an edge corrupted by noise
- Canny Edge Detector
 - ▶ Uses a mathematical model of the edge and the noise
 - ▶ Formalizes a performance criteria
 - ▶ Synthesizes the best filter

Canny's Criteria

- Detection – the important edges should not be missed and there should be no spurious responses
- Localization – the distance between the actual and located position of the edge should be minimal
- Single response – minimize multiple responses to a single edge
 - ▶ This is partly covered by the first criterion (since if there are two responses to a single edge, one of them should be considered as false)
 - ▶ This criterion solves the problem of an edge corrupted by noise
- Canny Edge Detector
 - ▶ Uses a mathematical model of the edge and the noise
 - ▶ Formalizes a performance criteria
 - ▶ Synthesizes the best filter

Canny's Criteria


- Detection – the important edges should not be missed and there should be no spurious responses
- Localization – the distance between the actual and located position of the edge should be minimal
- Single response – minimize multiple responses to a single edge
 - ▶ This is partly covered by the first criterion (since if there are two responses to a single edge, one of them should be considered as false)
 - ▶ This criterion solves the problem of an edge corrupted by noise
- Canny Edge Detector
 - ▶ Uses a mathematical model of the edge and the noise
 - ▶ Formalizes a performance criteria
 - ▶ Synthesizes the best filter

Performance Criteria

- Good detection
 - ▶ The filter must have a stronger response at the edge location ($x = 0$) than to noise
- Good localization
 - ▶ The filter response must be maximum very close to $x = 0$
- Low false positives
 - ▶ There should be only one maximum in a reasonable neighborhood of $x = 0$


Performance Criteria

- Good detection
 - ▶ The filter must have a stronger response at the edge location ($x = 0$) than to noise
- Good localization
 - ▶ The filter response must be maximum very close to $x = 0$
- Low false positives
 - ▶ There should be only one maximum in a reasonable neighborhood of $x = 0$




Performance Criteria

- Good detection
 - ▶ The filter must have a stronger response at the edge location ($x = 0$) than to noise
- Good localization
 - ▶ The filter response must be maximum very close to $x = 0$
- Low false positives
 - ▶ There should be only one maximum in a reasonable neighborhood of $x = 0$



Performance Criteria

- Good detection
 - ▶ The filter must have a stronger response at the edge location ($x = 0$) than to noise
- Good localization
 - ▶ The filter response must be maximum very close to $x = 0$
- Low false positives
 - ▶ There should be only one maximum in a reasonable neighborhood of $x = 0$



Canny Edge Detector

- Canny found a linear, continuous filter that maximized the three given criteria
 - ▶ There is no close-form solution for the optimal filter
 - ▶ However, it looks “very similar” to the first derivative of a Gaussian (DoG)
- Strictly speaking, Canny’s filter was derived in the context of 1-D profiles
 - ▶ To extend them to 2-D images we can run the filter at several different orientations in the image and detect edge elements at all orientations
- Canny Approximation
 - ▶ Smooth image with Gaussian
 - ▶ Compute the derivatives in x and y directions
 - ▶ Perform non-maximal suppression and sub-pixel interpolation using edge strength and direction values
 - ▶ Perform edge linking / hysteresis thresholding

Canny Edge Detector

- Canny found a linear, continuous filter that maximized the three given criteria
 - ▶ There is no close-form solution for the optimal filter
 - ▶ However, it looks “very similar” to the first derivative of a Gaussian (DoG)
- Strictly speaking, Canny’s filter was derived in the context of 1-D profiles
 - ▶ To extend them to 2-D images we can run the filter at several different orientations in the image and detect edge elements at all orientations
- Canny Approximation
 - ▶ Smooth image with Gaussian
 - ▶ Compute the derivatives in x and y directions
 - ▶ Perform non-maximal suppression and sub-pixel interpolation using edge strength and direction values
 - ▶ Perform edge linking / hysteresis thresholding

Canny Edge Detector

- Canny found a linear, continuous filter that maximized the three given criteria
 - ▶ There is no close-form solution for the optimal filter
 - ▶ However, it looks “very similar” to the first derivative of a Gaussian (DoG)
- Strictly speaking, Canny’s filter was derived in the context of 1-D profiles
 - ▶ To extend them to 2-D images we can run the filter at several different orientations in the image and detect edge elements at all orientations
- Canny Approximation
 - ▶ Smooth image with Gaussian
 - ▶ Compute the derivatives in x and y directions
 - ▶ Perform non-maximal suppression and sub-pixel interpolation using edge strength and direction values
 - ▶ Perform edge linking / hysteresis thresholding

1-D Gaussian

- 1-D Gaussian with zero mean is given by



$$g(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}}$$

- Ignore the scale factor,

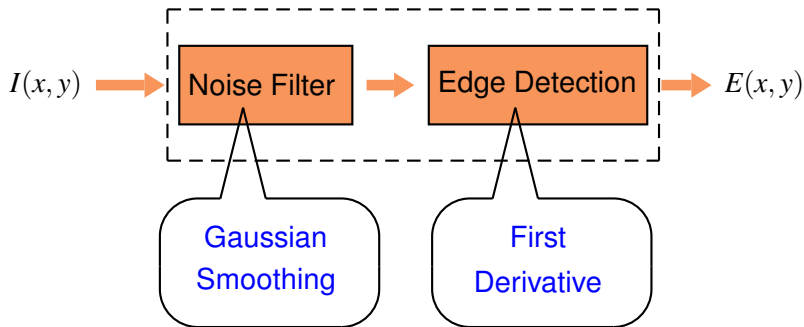
$$g(x) = e^{-\frac{x^2}{2\sigma^2}}$$

- The first derivative of Gaussian is given by

$$g'(x) = -\frac{1}{2\sigma^2} 2xe^{-\frac{x^2}{2\sigma^2}} = -\frac{x^2}{\sigma^2} e^{-\frac{x^2}{2\sigma^2}}$$



Another Interpretation



$$E(x) = \frac{d(I(x) * G(x))}{dx} = I(x) * \frac{dG(x)}{dx}$$

Canny Edge Detector

- 1-D
 - First derivative
 - $E(x) = \frac{d(I(x) * G(x))}{dx}$
 - Absolute value
 - $|E(x)| \geq Th$
- 2-D
 - Gradient vector
 - $E(x, y) = \nabla(I(x, y) * G(x, y))$
 - Magnitude
 - $|E(x)| \geq Th$

CANNY_ENHANCER

- The input is image I ; G is a zero mean Gaussian filter with standard derivation σ
 - ▶ $J = I * G$ (smoothing)
 - ▶ For each pixel (i, j) : (edge enhancement)
 - ★ Compute the image gradient: $\nabla J(i, j) = (J_x(i, j), J_y(i, j))$
 - ★ Estimate edge strength: $e_s(i, j) = \sqrt{J_x^2(i, j) + J_y^2(i, j)}$
 - ★ Estimate edge orientation: $e_o = \tan^{-1} \frac{J_y(i, j)}{J_x(i, j)}$
 - The output are images E_s and E_o (edge strength and edge orientation)
-
- The output image E_s has the magnitudes of the smoothed gradient
 - σ determines the amount of smoothing
 - E_s has large values at edges


CANNY_ENHANCER

- The input is image I ; G is a zero mean Gaussian filter with standard derivation σ
 - ▶ $J = I * G$ (smoothing)
 - ▶ For each pixel (i, j) : (edge enhancement)
 - ★ Compute the image gradient: $\nabla J(i, j) = (J_x(i, j), J_y(i, j))$
 - ★ Estimate edge strength: $e_s(i, j) = \sqrt{J_x^2(i, j) + J_y^2(i, j)}$
 - ★ Estimate edge orientation: $e_o = \tan^{-1} \frac{J_x(i, j)}{J_y(i, j)}$
- The output are images E_s and E_o (edge strength and edge orientation)
- The output image E_s has the magnitudes of the smoothed gradient
- σ determines the amount of smoothing
- E_s has large values at edges


Edge Detection

- E_s has large values at edges:

- ▶ Find local maxima



- ▶ But it also may have wide ridges around the local maxima (large value around the edges)



Non-Maximal Suppression

- One approach to overcoming the problem of edge thickening is to explicitly look for maxima of the response to the edge enhancement filter
- Sub-pixel Localization
 - ▶ One can try to further localize the position of the edge within a pixel by analyzing the response to the edge enhancement filter
 - ▶ One common approach is to fit a “quadratic polynomial” to the filter response in the region of a maxima and compute the true maximum
 - ▶ Let $y(x) = ax^2 + bx + c$ and perform interpolation to find max or min

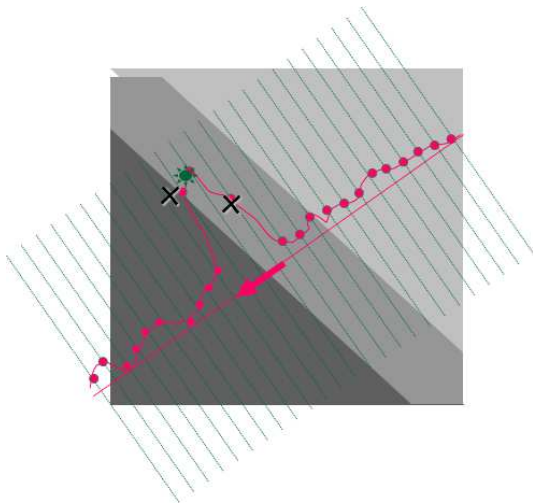
Non-Maximal Suppression

- One approach to overcoming the problem of edge thickening is to explicitly look for maxima of the response to the edge enhancement filter
- Sub-pixel Localization
 - ▶ One can try to further localize the position of the edge within a pixel by analyzing the response to the edge enhancement filter
 - ▶ One common approach is to fit a “quadratic polynomial” to the filter response in the region of a maxima and compute the true maximum
 - ▶ Let $y(x) = ax^2 + bx + c$ and perform interpolation to find max or min

NONMAX_SUPPRESSION

- The inputs are E_s and E_o (outputs of CANNY_ENHANCER)
- Consider 4 directions $D = \{0, 45, 90, 135 \text{ degrees}\}$ with respect to horizontal axis image reference frame
- For each pixel (i, j) do:
 - ▶ Find the direction $d \in D$ such that $d \approx E_o(i, j)$ (normal to the edge)
 - ▶ If $E_s(i, j)$ smaller than at least one of its neighbors along d
 - ★ $I_N(i, j) = 0$ (suppression)
 - ★ Otherwise, $I_N(i, j) = E_s(i, j)$
- The output is the thinned edge image I_N

Graphical Interpretation



Edge Linking

- The output of the edge enhancement stage is a binary array indicating the locations of edgels (edge elements) in the image
- The edge linking stage attempts to group these discrete elements into chains much like stringing pearls
- Problems in edge linking:
 - ▶ Edges can be broken because of low contrast
 - ▶ Junctions can cause major problems since the edge enhancement procedure tends to fail in these situations and the linker can become confused

Thresholding


- Edges are found by thresholding the output of NONMAX_SUPPRESSION
- If the threshold is too high: Very few (none) edges
 - ▶ High misdetections
 - ▶ Many gaps
- If the threshold is too low: Too many (all pixels) edges
 - ▶ High false positives
 - ▶ Many extra edges

Hysteresis Thresholding

- Canny proposed an approach to dealing with broken edge chains in the linking phase
- The idea is to maintain two thresholds on edge strength one for starting a chain and a lower one for use during linking
- In this way the linker will work well even when a chain has some low contrast sections

Solution


- Hysteresis Thresholding
- “Strong edges” reinforce adjacent “weak edges”



HYSTERESIS_THRESH

- Inputs:
 - ▶ I_N (output of NONMAX_SUPPRESSION)
 - ▶ E_o (output of CANNY_ENHANCER)
 - ▶ Thresholds L and H
- Scanning all edge points in I_N in a fixed order:
 - ▶ Locate the next unvisited pixel such that $I_N(i,j) > H$
 - ▶ Starting from $I_N(i,j)$, follow the chains of connected local maxima, in both directions perpendicular to the edge normal, as long as $I_N > L$; Mark all visited points, and save the location of the contour points
- Output:
 - ▶ A set of lists describing the contours

Hysteresis Thresholding



Algorithm: Canny Edge Detector

- Convolve an image f with a Gaussian of scale σ
- Estimate local edge normal directions n for each pixel in the image
- Find the location of the edges (non-maximal suppression)
- Compute the magnitude of the edge
- Threshold edges in the image with hysteresis to eliminate spurious responses
- Repeat steps 1. through 5. for ascending values of the standard deviation σ
- Aggregate the final information about edges at multiple scale using the “feature synthesis” approach

Algorithm: Canny Edge Detector

- Convolve an image f with a Gaussian of scale σ
- Estimate local edge normal directions n for each pixel in the image
- Find the location of the edges (non-maximal suppression)
- Compute the magnitude of the edge
- Threshold edges in the image with hysteresis to eliminate spurious responses
- Repeat steps 1. through 5. for ascending values of the standard deviation σ
- Aggregate the final information about edges at multiple scale using the “feature synthesis” approach

Algorithm: Canny Edge Detector

- Convolve an image f with a Gaussian of scale σ
- Estimate local edge normal directions n for each pixel in the image
- Find the location of the edges (non-maximal suppression)
- Compute the magnitude of the edge
- Threshold edges in the image with hysteresis to eliminate spurious responses
- Repeat steps 1. through 5. for ascending values of the standard deviation σ
- Aggregate the final information about edges at multiple scale using the “feature synthesis” approach

Algorithm: Canny Edge Detector

- Convolve an image f with a Gaussian of scale σ
- Estimate local edge normal directions n for each pixel in the image
- Find the location of the edges (non-maximal suppression)
- Compute the magnitude of the edge
- Threshold edges in the image with hysteresis to eliminate spurious responses
- Repeat steps 1. through 5. for ascending values of the standard deviation σ
- Aggregate the final information about edges at multiple scale using the “feature synthesis” approach

Algorithm: Canny Edge Detector

- Convolve an image f with a Gaussian of scale σ
- Estimate local edge normal directions n for each pixel in the image
- Find the location of the edges (non-maximal suppression)
- Compute the magnitude of the edge
- Threshold edges in the image with hysteresis to eliminate spurious responses
- Repeat steps 1. through 5. for ascending values of the standard deviation σ
- Aggregate the final information about edges at multiple scale using the “feature synthesis” approach

Algorithm: Canny Edge Detector

- Convolve an image f with a Gaussian of scale σ
- Estimate local edge normal directions n for each pixel in the image
- Find the location of the edges (non-maximal suppression)
- Compute the magnitude of the edge
- Threshold edges in the image with hysteresis to eliminate spurious responses
- Repeat steps 1. through 5. for ascending values of the standard deviation σ
- Aggregate the final information about edges at multiple scale using the “feature synthesis” approach

Algorithm: Canny Edge Detector

- Convolve an image f with a Gaussian of scale σ
- Estimate local edge normal directions n for each pixel in the image
- Find the location of the edges (non-maximal suppression)
- Compute the magnitude of the edge
- Threshold edges in the image with hysteresis to eliminate spurious responses
- Repeat steps 1. through 5. for ascending values of the standard deviation σ
- Aggregate the final information about edges at multiple scale using the “feature synthesis” approach

Example



Threshold = 0.1, $\sigma = 1$



Threshold = 0.1, $\sigma = 2.8$



Threshold = 0.3, $\sigma = 1$



Threshold = 0.3, $\sigma = 2.8$

Disclaimer

- No edge detection scheme is going to work perfectly in all cases
- This is due to the fact that our notion of what constitutes a salient edge in the image is actually somewhat subtle

Detecting Corners

- We are often interested in detecting point features in an image
- These features are usually defined as regions in the image where there is significant edge strength in two or more directions
- They can be used for
 - ▶ Object tracking
 - ▶ 3D triangulation (stereo)
 - ▶ Object recognition
- Need two strong edges
- If E_x and E_y denote the gradients of the image in the x and y directions, then the behavior of the gradients in a region around a point can be obtained by considering the following matrix

$$C = \sum \begin{pmatrix} E_x \\ E_y \end{pmatrix} \begin{pmatrix} E_x & E_y \end{pmatrix} = \sum \begin{pmatrix} E_x^2 & E_x E_y \\ E_x E_y & E_y^2 \end{pmatrix}$$

Detecting Corners

- We are often interested in detecting point features in an image
- These features are usually defined as regions in the image where there is significant edge strength in two or more directions
- They can be used for
 - ▶ Object tracking
 - ▶ 3D triangulation (stereo)
 - ▶ Object recognition
- Need two strong edges
- If E_x and E_y denote the gradients of the image in the x and y directions, then the behavior of the gradients in a region around a point can be obtained by considering the following matrix

$$C = \sum \begin{pmatrix} E_x \\ E_y \end{pmatrix} \begin{pmatrix} E_x & E_y \end{pmatrix} = \sum \begin{pmatrix} E_x^2 & E_x E_y \\ E_x E_y & E_y^2 \end{pmatrix}$$

Examining The Matrix

- One way to decide on the presence of a corner is to look at the eigenvalues of the 2×2 matrix C
 - ▶ If the area is a region of “constant intensity” we would expect both eigenvalues to be small (or zero)
 - ▶ If it contains a edge we expect one large eigenvalue and one small one
 - ▶ If it contains edges at two or more orientations we expect two large eigenvalues
- If $\min(\lambda_1, \lambda_2) > T$, then there is a corner!

Finding Corner

- One approach to finding corners is to find locations where the smaller eigenvalue is greater than some threshold
- We could also consider the ratio of the two eigenvalues
- Issues:
 - ▶ Localization – It can be difficult to precisely localize the corner in the intensity image
 - ▶ Modeling – It can be helpful to have a model of the corners you are trying to find in order to detect and localize them more systematically

Finding Corner


- One approach to finding corners is to find locations where the smaller eigenvalue is greater than some threshold
- We could also consider the ratio of the two eigenvalues
- Issues:
 - ▶ Localization – It can be difficult to precisely localize the corner in the intensity image
 - ▶ Modeling – It can be helpful to have a model of the corners you are trying to find in order to detect and localize them more systematically

Corner Detection



- Compute image gradient
- For each $m \times m$ neighborhood, compute matrix C
- If smaller eigenvalue λ_2 is greater than threshold τ , record a corner
- Non-maximum suppression: only keep strongest corner in each $m \times m$ window

Corner Detection Results

- Checkerboard with noise



Corner Detection Results



Histogram of λ_2 (smaller eigenvalues)

Detecting Lines

- The difference between line detection and edge detection:
 - ▶ Edges = local
 - ▶ Lines = non-local
- Line detection usually performed on the output of an edge detector
- Several different approaches:
 - ▶ For each possible line, check whether the line is present: “brute force”
 - ▶ Given detected edges, record lines to which they might belong: “Hough transform + voting”
 - ▶ Given guess for approximate location of a line, refine that guess: “fitting”
- Second method (Hough transform) is efficient for finding unknown lines, but not always accurate

Detecting Lines


- The difference between line detection and edge detection:
 - ▶ Edges = local
 - ▶ Lines = non-local
- Line detection usually performed on the output of an edge detector
- Several different approaches:
 - ▶ For each possible line, check whether the line is present: “brute force”
 - ▶ Given detected edges, record lines to which they might belong: “Hough transform + voting”
 - ▶ Given guess for approximate location of a line, refine that guess: “fitting”
- Second method (Hough transform) is efficient for finding unknown lines, but not always accurate

Detecting Lines

- The difference between line detection and edge detection:
 - ▶ Edges = local
 - ▶ Lines = non-local
- Line detection usually performed on the output of an edge detector
- Several different approaches:
 - ▶ For each possible line, check whether the line is present: “brute force”
 - ▶ Given detected edges, record lines to which they might belong: “Hough transform + voting”
 - ▶ Given guess for approximate location of a line, refine that guess: “fitting”
- Second method (Hough transform) is efficient for finding unknown lines, but not always accurate

Line Detection

Mathematical model of a line:



$$y_1 = mx_1 + n$$


$$y_2 = mx_2 + n$$

$$\vdots$$

$$y_N = mx_N + n$$

Line Detection

Mathematical model of a line:



$$y_1 = mx_1 + n$$


$$y_2 = mx_2 + n$$

$$\vdots$$

$$y_N = mx_N + n$$

Line Detection

Mathematical model of a line:



$$y_1 = mx_1 + n$$


$$y_2 = mx_2 + n$$

\vdots

$$y_N = mx_N + n$$

Line Detection

Mathematical model of a line:



$$y_1 = mx_1 + n$$


$$y_2 = mx_2 + n$$

$$\vdots$$

$$y_N = mx_N + n$$

Line Detection

Mathematical model of a line:



$$y_1 = mx_1 + n$$

$$y_2 = mx_2 + n$$

$$\vdots$$

$$y_N = mx_N + n$$

Image and Parameter Spaces





Image Space

$$\begin{array}{l} y_1 = mx_1 + n \\ y_2 = mx_2 + n \\ \vdots \\ y_N = mx_N + n \end{array}$$

$$\begin{array}{l} y_1 = m'x_1 + n' \\ y_2 = m'x_2 + n' \\ \vdots \\ y_N = m'x_N + n' \end{array}$$



Parameter Space

Line in Image Space \sim Point in Parameter Space

Image and Parameter Spaces





Image Space

$$\begin{aligned}y_1 &= mx_1 + n \\y_2 &= mx_2 + n \\&\vdots \\y_N &= mx_N + n\end{aligned}$$

$$\begin{aligned}y_1 &= m'x_1 + n' \\y_2 &= m'x_2 + n' \\&\vdots \\y_N &= m'x_N + n'\end{aligned}$$



Parameter Space

Line in Image Space \sim Point in Parameter Space

Image and Parameter Spaces




Image Space

$$y_1 = mx_1 + n$$

$$y_2 = mx_2 + n$$

$$\vdots$$


$$y_N = mx_N + n$$

$$y_1 = m'x_1 + n'$$

$$y_2 = m'x_2 + n'$$

$$\vdots$$

$$y_N = m'x_N + n'$$



Parameter Space

Line in Image Space \sim Point in Parameter Space

Image and Parameter Spaces




Image Space

$$y_1 = mx_1 + n$$

$$y_2 = mx_2 + n$$

$$\vdots$$


$$y_N = mx_N + n$$

$$y_1 = m'x_1 + n'$$

$$y_2 = m'x_2 + n'$$

$$\vdots$$

$$y_N = m'x_N + n'$$



Parameter Space

Line in Image Space \sim Point in Parameter Space

Image and Parameter Spaces




Image Space

$$y_1 = mx_1 + n$$

$$y_2 = mx_2 + n$$

$$\vdots$$


$$y_N = mx_N + n$$

$$y_1 = m'x_1 + n'$$

$$y_2 = m'x_2 + n'$$

$$\vdots$$

$$y_N = m'x_N + n'$$



Parameter Space

Line in Image Space \sim Point in Parameter Space

Looking at it Backwards ...


Image Space

Fix (m, n) , vary (x, y) – Line

Fix (x_1, y_1) , vary (m, n) – Lines through a Point

$$y = mx + n$$

$$y_1 = mx_1 + n$$



Looking at it Backwards ...

Parameter Space

The line $y_1 = mx_1 + n$ can be re-written as:

Fix $(-x_1, y_1)$, vary (m, n) – Line

$$n = -x_1 m + y_1$$

$$n = -x_1 m + y_1$$




Image & Parameter Spaces

Image Space

- Lines
- Points
- Collinear points

Parameter Space

- Points
- Lines
- Intersecting lines

This is called duality!

Hough Transform




- General idea: transform from image coordinates to parameter space of features
 - ▶ Map a difficult pattern problem into a simple peak detection problem
 - ▶ Need parameterized model of features
 - ▶ For each pixel, determine all parameter values that might have given rise to that pixel; vote
 - ▶ At end, look for peaks in parameter space
- This approach is a voting scheme based on accumulating evidence in a parameter space

Hough Transform

- General idea: transform from image coordinates to parameter space of features
 - ▶ Map a difficult pattern problem into a simple peak detection problem
 - ▶ Need parameterized model of features
 - ▶ For each pixel, determine all parameter values that might have given rise to that pixel; vote
 - ▶ At end, look for peaks in parameter space
- This approach is a voting scheme based on accumulating evidence in a parameter space

Hough Transform for Lines

- Each input measurement indicates its contribution to a globally consistent solution
- Here this problem is under constrained
 - ▶ Generic line: $y = ax + b$
 - ▶ Parameters: a and b



Hough Transform for Lines

- Given an edge point, there is an infinite number of lines passing through it (vary m and n)
- These lines can be represented as a line in parameter space





Image Space



Parameter Space

Hough Transform for Lines

- Given a set of collinear edge points, each of them have associated a line in parameter spaces
- These lines intersect at the point (m, n) corresponding to the parameters of the line in the image space





Image Space



Parameter Space

Hough Transform Technique


- At each point of the (discrete) parameter space, count how many lines pass through it
 - ▶ Use an array of counters
 - ▶ Can be thought as a “parameter image”
- The higher the count, the more edges are collinear in the image space
 - ▶ Find a peak in the counter array
 - ▶ This is a “bright” point in the parameter image
 - ▶ It can be found by thresholding
- Practical Issues
 - ▶ The slope of the line is $-\infty < m < \infty$
 - ★ The parameter space is infinite
 - ▶ The representation $y = mx + n$ does not express lines of the form $x = k$

Hough Transform Technique

- At each point of the (discrete) parameter space, count how many lines pass through it
 - ▶ Use an array of counters
 - ▶ Can be thought as a “parameter image”
- The higher the count, the more edges are collinear in the image space
 - ▶ Find a peak in the counter array
 - ▶ This is a “bright” point in the parameter image
 - ▶ It can be found by thresholding
- Practical Issues
 - ▶ The slope of the line is $-\infty < m < \infty$
 - ★ The parameter space is infinite
 - ▶ The representation $y = mx + n$ does not express lines of the form $x = k$

Solution

- Use the “normal” equation of a line:



$$\rho = x \cos \theta + y \sin \theta$$

θ is the line orientation

ρ is the distance between the origin and the line

New Parameter Space

- Use the parameter space (ρ, θ)
- The new space is finite
 - ▶ $0 < \rho < D$, where D is the image diagonal
 - ▶ $0 < \theta < 2\pi$
- The new space can represent all lines
 - ▶ $y = k$ is represented with $\rho = k, \theta = 90^\circ$
 - ▶ $x = k$ is represented with $\rho = k, \theta = 0^\circ$
- A point in image space is now represented as a sinusoid
 - ▶ $\rho = x \cos \theta + y \sin \theta$

New Parameter Space

- Use the parameter space (ρ, θ)
- The new space is finite
 - ▶ $0 < \rho < D$, where D is the image diagonal
 - ▶ $0 < \theta < 2\pi$
- The new space can represent all lines
 - ▶ $y = k$ is represented with $\rho = k, \theta = 90^\circ$
 - ▶ $x = k$ is represented with $\rho = k, \theta = 0^\circ$
- A point in image space is now represented as a sinusoid
 - ▶ $\rho = x \cos \theta + y \sin \theta$



Hough Transform Algorithm

- Input is an edge image ($E(i,j) = 1$ for edgels)
 - ▶ Discretize θ and ρ in increments of θ_d and ρ_d
 - ▶ Let $A(R, T)$ be an array of integer accumulators, initialized to 0
 - ▶ For each pixel $E(i,j) = 1$ and $h = 1, 2, \dots, T$ do
 - ★ $\rho = i \cos(h\theta_d) + j \sin(h\theta_d)$
 - ★ Find closest integer k of the element of ρ_d , corresponding to ρ
 - ★ Increment counter $A(h, k)$ by one
 - ▶ Find all local maxima in $A(R, T) > threshold$
- Output is a set of pairs (ρ_d, θ_d) describing the lines detected in E in polar form


Hough Transform Speed Up

- If we know the orientation of the edge – usually available from the edge detection step
 - ▶ We fix θ in the parameter space and increment only one counter!
 - ▶ We can allow for orientation uncertainty by incrementing a few counters around the “nominal” counter

Example





Example



Example



Example



Active or Deformable Contours

- How to fit a curve of arbitrary shape to a set of image edge points? (restricted to closed contours only)
- General closed curves can be represented by snake (also called active contour or deformable contour)
- Deformable models represent
 - ▶ Class of objects of differing shape (bananas)
 - ▶ Objects which change shape (such as lips)
- Deformable models may be
 - ▶ 3D surface, (a balloon squeezed out of shape)
 - ▶ 3D space curves, which we bend to form figures
 - ▶ 2D contours, e.g. the Snake

Active or Deformable Contours

- How to fit a curve of arbitrary shape to a set of image edge points? (restricted to closed contours only)
- General closed curves can be represented by snake (also called active contour or deformable contour)
- Deformable models represent
 - ▶ Class of objects of differing shape (bananas)
 - ▶ Objects which change shape (such as lips)
- Deformable models may be
 - ▶ 3D surface, (a balloon squeezed out of shape)
 - ▶ 3D space curves, which we bend to form figures
 - ▶ 2D contours, e.g. the Snake

Deformable Contours

- Goal
 - ▶ Start with image and initial closed curve
 - ▶ Evolve curve to lie along “important” feature: Edges, Corners, Detected features, User input
- The concept of a snake applied to computer vision
 - ▶ It is an elastic band of arbitrary shape
 - ▶ It is sensitive to the image gradient
 - ▶ Initially it is located near the image contour of interest
 - ▶ It can wiggle in the image
 - ▶ It is represented as a necklace of points
 - ▶ It is then attracted towards the target contour by forces depending on the intensity gradient

Deformable Contours

- Goal
 - ▶ Start with image and initial closed curve
 - ▶ Evolve curve to lie along “important” feature: Edges, Corners, Detected features, User input
- The concept of a snake applied to computer vision
 - ▶ It is an elastic band of arbitrary shape
 - ▶ It is sensitive to the image gradient
 - ▶ Initially it is located near the image contour of interest
 - ▶ It can wiggle in the image
 - ▶ It is represented as a necklace of points
 - ▶ It is then attracted towards the target contour by forces depending on the intensity gradient

Active Contour Models – Snakes

- Introduced by Kass, Witkin, and Terzopoulos
- Framework: energy minimization
 - ▶ Bending and stretching curve = more energy
 - ▶ Good features = less energy
 - ▶ Curve evolves to minimize energy
- The key idea of deformable contour
 - ▶ To associate an energy functional to each possible contour shape such that the image contour to be detected corresponds to a minimum of functional
 - ▶ The snake is applied to the intensity image
 - ▶ Other curve fitting algorithms are applied to edge points

- User-Visible Options

- ▶ Initialization: user-specified, automatic
- ▶ Curve properties: continuity, smoothness
- ▶ Image features: intensity, edges, corners, ...
- ▶ Other forces: hard constraints, springs, attractors, ...
- ▶ Scale: local, multi-resolution, global

- Behind-the-Scenes Options

- ▶ Framework: energy minimization, forces acting on curve
- ▶ Curve representation: ideal curve, sampled, spline, implicit function
- ▶ Evolution method: calculus of variations, numerical differential equations, local search

- User-Visible Options


- ▶ Initialization: user-specified, automatic
- ▶ Curve properties: continuity, smoothness
- ▶ Image features: intensity, edges, corners, ...
- ▶ Other forces: hard constraints, springs, attractors, ...
- ▶ Scale: local, multi-resolution, global

- Behind-the-Scenes Options

- ▶ Framework: energy minimization, forces acting on curve
- ▶ Curve representation: ideal curve, sampled, spline, implicit function
- ▶ Evolution method: calculus of variations, numerical differential equations, local search

Main Idea

- “Drop” a snake
- Let the snake “wiggle” attracted by image gradient, until it glues itself against a contour



Introductory Demo

Active Contour Model – Snakes

- Active contour models may be used in image segmentation and understanding, and are also suitable for analysis of dynamic image data or 3D image data
- It is defined as an energy-minimizing spline – the snake's energy depends on its shape and location within the image
- Local minima of this energy then correspond to desired image properties
- The snake is active, always minimizing its energy functional, therefore exhibiting dynamic behavior
- The idea behind deformable contours is to find a contour $c(s)$ which best approximates the perimeter of an object
- The approach is to construct an energy functional which measures the appropriateness of a contour and to optimize this functional with respect to the contour parameters

Active Contour Model – Snakes

- Active contour models may be used in image segmentation and understanding, and are also suitable for analysis of dynamic image data or 3D image data
- It is defined as an energy-minimizing spline – the snake's energy depends on its shape and location within the image
- Local minima of this energy then correspond to desired image properties
- The snake is active, always minimizing its energy functional, therefore exhibiting dynamic behavior
- The idea behind deformable contours is to find a contour $c(s)$ which best approximates the perimeter of an object
- The approach is to construct an energy functional which measures the appropriateness of a contour and to optimize this functional with respect to the contour parameters

Active Contour Model – Snakes

- Active contour models may be used in image segmentation and understanding, and are also suitable for analysis of dynamic image data or 3D image data
- It is defined as an energy-minimizing spline – the snake's energy depends on its shape and location within the image
- Local minima of this energy then correspond to desired image properties
- The snake is active, always minimizing its energy functional, therefore exhibiting dynamic behavior
- The idea behind deformable contours is to find a contour $c(s)$ which best approximates the perimeter of an object
- The approach is to construct an energy functional which measures the appropriateness of a contour and to optimize this functional with respect to the contour parameters

Active Contour Model – Snakes

- Active contour models may be used in image segmentation and understanding, and are also suitable for analysis of dynamic image data or 3D image data
- It is defined as an energy-minimizing spline – the snake's energy depends on its shape and location within the image
- Local minima of this energy then correspond to desired image properties
- The snake is active, always minimizing its energy functional, therefore exhibiting dynamic behavior
- The idea behind deformable contours is to find a contour $c(s)$ which best approximates the perimeter of an object
- The approach is to construct an energy functional which measures the appropriateness of a contour and to optimize this functional with respect to the contour parameters

Active Contour Model – Snakes

- Active contour models may be used in image segmentation and understanding, and are also suitable for analysis of dynamic image data or 3D image data
- It is defined as an energy-minimizing spline – the snake's energy depends on its shape and location within the image
- Local minima of this energy then correspond to desired image properties
- The snake is active, always minimizing its energy functional, therefore exhibiting dynamic behavior
- The idea behind deformable contours is to find a contour $c(s)$ which best approximates the perimeter of an object
- The approach is to construct an energy functional which measures the appropriateness of a contour and to optimize this functional with respect to the contour parameters

Active Contour Model – Snakes

- Active contour models may be used in image segmentation and understanding, and are also suitable for analysis of dynamic image data or 3D image data
- It is defined as an energy-minimizing spline – the snake's energy depends on its shape and location within the image
- Local minima of this energy then correspond to desired image properties
- The snake is active, always minimizing its energy functional, therefore exhibiting dynamic behavior
- The idea behind deformable contours is to find a contour $c(s)$ which best approximates the perimeter of an object
- The approach is to construct an energy functional which measures the appropriateness of a contour and to optimize this functional with respect to the contour parameters

Energy Functional

- Associate to each possible shape and location of the snake a value E
 - ▶ Values should be such that the image contour to be detected has the minimum value
 - ▶ E is called the energy of the snake
- Keep wiggling the snake towards smaller value
- We need a function that given a snake state, associates to it an energy value
- The function should be designed so that the snake moves towards the contour that we are seeking!

Energy Functional

- Associate to each possible shape and location of the snake a value E
 - ▶ Values should be such that the image contour to be detected has the minimum value
 - ▶ E is called the energy of the snake
- Keep wiggling the snake towards smaller value
- We need a function that given a snake state, associates to it an energy value
- The function should be designed so that the snake moves towards the contour that we are seeking!

What Moves The Snake

- Forces moving the snake (External)

- ▶ It needs to be attracted to contours:
 - ★ Edge pixels “pull” the snake points
 - ★ The stronger the edge, the stronger the pull
 - ★ The force is proportional to $|\nabla|$

- Forces preserving the snake (Internal)

- ▶ The snake should not break apart!
 - ★ Points on the snake must stay close to each other
 - ★ Each point on the snake pulls its neighbors
 - ★ The farther the neighbor, the stronger the force
 - ★ The force is proportional to the distance $|\mathbf{P}_i - \mathbf{P}_{i-1}|$
- ▶ The snake should avoid “oscillations”
 - ★ Penalize high curvature
 - ★ Force proportional to snake curvature

What Moves The Snake

- Forces moving the snake (External)
 - ▶ It needs to be attracted to contours:
 - ★ Edge pixels “pull” the snake points
 - ★ The stronger the edge, the stronger the pull
 - ★ The force is proportional to $|\nabla|$
- Forces preserving the snake (Internal)
 - ▶ The snake should not break apart!
 - ★ Points on the snake must stay close to each other
 - ★ Each point on the snake pulls its neighbors
 - ★ The farther the neighbor, the stronger the force
 - ★ The force is proportional to the distance $|\mathbf{P}_i - \mathbf{P}_{i-1}|$
 - ▶ The snake should avoid “oscillations”
 - ★ Penalize high curvature
 - ★ Force proportional to snake curvature

Energy Functional

- The minimized energy functional is a weighted combination of internal and external forces
 - ▶ Internal forces – emanate from the shape of the snake
 - ▶ External forces – come from the image and/or from high-level image understanding processes
- The snake is defined parametrically as $\mathbf{v}(s) = (x(s), y(s))$, where $x(s), y(s)$ are x, y coordinates along the contour and $s \in [0, 1]$
- The energy functional to be minimized may be written as

$$E_{snake}^* = \int_0^1 E_{snake}(\mathbf{v}(s)) ds = \int_0^1 \{[E_{int}(\mathbf{v}(s))] + [E_{image}(\mathbf{v}(s))] + [E_{con}(\mathbf{v}(s))]\}$$

- ▶ E_{int} – the internal energy of the spline due to bending
- ▶ E_{image} – image forces
- ▶ E_{con} – external constraint

Energy Functional

- The minimized energy functional is a weighted combination of internal and external forces
 - ▶ Internal forces – emanate from the shape of the snake
 - ▶ External forces – come from the image and/or from high-level image understanding processes
- The snake is defined parametrically as $\mathbf{v}(s) = (x(s), y(s))$, where $x(s), y(s)$ are x, y coordinates along the contour and $s \in [0, 1]$
- The energy functional to be minimized may be written as

$$E_{snake}^* = \int_0^1 E_{snake}(\mathbf{v}(s)) ds = \int_0^1 \{[E_{int}(\mathbf{v}(s))] + [E_{image}(\mathbf{v}(s))] + [E_{con}(\mathbf{v}(s))]\}$$

- ▶ E_{int} – the internal energy of the spline due to bending
- ▶ E_{image} – image forces
- ▶ E_{con} – external constraint

Internal Energy

- The internal spline energy can be written as

$$E_{int} = \alpha(s) \left| \frac{d\mathbf{v}}{ds} \right|^2 + \beta(s) \left| \frac{d^2\mathbf{v}}{ds^2} \right|^2$$

where $\alpha(s), \beta(s)$ specify the elasticity and stiffness of the snake

- First term is “membrane” term – minimum energy when curve minimizes length (“soap bubble”)
- Second term is “thin plate” term – minimum energy when curve is smooth
- Control α and β to vary between extremes
- Set β to 0 at a point to allow corner (2nd-order discontinuous)

Image Energy

- The second term of the energy integral is derived from the image data over which the snake lies
- Variety of terms gives different effects
- For example, a weighted combination of three different functionals is presented which attracts the snake to lines, edges and terminations:

$$E_{image} = w_{int}E_{line} + w_{edge}E_{edge} + w_{term}E_{term}$$

- The line-based functional $E_{line} = f(x, y)$
- The edge-based functional $E_{edge} = -|\nabla f(x, y)|^2$ attracts the snake to contours with large image gradients (location of strong edges)
- Line terminations and corners may influence the snake using a weighted energy functional $E_{term} = \frac{\partial \phi}{\partial \mathbf{n}_R}$ where $\phi(x, y)$ denotes the gradient direction along the spline, etc.

Constraint Forces

- The third term of the integral comes from external constraints imposed either by a user or some other high-level process which may force the snake toward or away from particular features
- If the snake is near to some desired feature, the energy minimization will pull the snake the rest of the way
- If the snake settles in a local energy minimum that a high-level process determines as incorrect, an area of energy peak may be made at this location to force the snake away to a different local minimum
- Spring: $E_{con} = k|\mathbf{v} - \mathbf{x}|^2$
- Repulsion: $E_{con} = \frac{k}{|\mathbf{v} - \mathbf{x}|^2}$

Minimization


- A contour is defined to lie in the position in which the snake reaches a local energy minimum
- The functional to be minimized is

$$E_{snake}^* = \int_0^1 E_{snake}[\mathbf{v}(s)] ds$$


- From the calculus of variations, the Euler-Lagrange condition states that the spline $\mathbf{v}(s)$ which minimizes E_{snake}^* must satisfy

$$-\frac{d^2}{ds^2} \left(\frac{\partial E}{\partial \left(\frac{d^2 x}{ds^2} \right)} + \frac{\partial E}{\partial \left(\frac{d^2 y}{ds^2} \right)} \right) + \frac{d}{ds} E_{\mathbf{v}_s} - E_{\mathbf{v}} = 0$$

Corpus Callosum



Corpus Callosum



Evolving Curve

- Computing forces on \mathbf{v} that locally minimize energy gives differential equation for \mathbf{v}
- Discretize \mathbf{v} : samples (x_i, y_i)
- Approximate derivatives using finite differences
- Numerical solver iteratively converges to minimum
- Write equations directly in terms of forces, not energy
- Implicit equation solver
- Search neighborhood of each (x_i, y_i) for pixel that minimizes energy
- Exact solution: calculus of variations

Some Comments

● Advantage of Snakes

- ▶ Easy to manipulate (intuitive)
- ▶ Sensitive to image scale by Gaussian smoothing in the image energy function
- ▶ Insensitive to noise and other ambiguities in the images
- ▶ They can be used to track dynamic objects in temporal as well as the spatial dimensions

● Disadvantage of Snakes

- ▶ Often get stuck in local minima states
 - ★ Overcome by simulated annealing techniques at the expense of longer computation times
- ▶ Often overlook minute features in the process of minimizing the energy over the entire path of their contours
- ▶ Their accuracy is governed by the convergence criteria used in the energy minimization technique – Higher accuracy requires tighter convergence criteria and longer computation

Some Comments


- Advantage of Snakes

- ▶ Easy to manipulate (intuitive)
- ▶ Sensitive to image scale by Gaussian smoothing in the image energy function
- ▶ Insensitive to noise and other ambiguities in the images
- ▶ They can be used to track dynamic objects in temporal as well as the spatial dimensions

- Disadvantage of Snakes

- ▶ Often get stuck in local minima states
 - ★ Overcome by simulated annealing techniques at the expense of longer computation times
- ▶ Often overlook minute features in the process of minimizing the energy over the entire path of their contours
- ▶ Their accuracy is governed by the convergence criteria used in the energy minimization technique – Higher accuracy requires tighter convergence criteria and longer computation

Brain Cortex Segmentation



Add energy term for constant-color regions of a single color

Scale

- In the simplest snakes algorithm, image features only attract locally
- Greater region of attraction: smooth image
 - ▶ Curve might not follow high-frequency detail
- Multi-resolution processing
- Looking for global minimum vs. local minima
 - ▶ Start with smoothed image to attract curve
 - ▶ Finish with unsmoothed image to get details

Scale

- In the simplest snakes algorithm, image features only attract locally
- Greater region of attraction: smooth image
 - ▶ Curve might not follow high-frequency detail
- Multi-resolution processing
- Looking for global minimum vs. local minima
 - ▶ Start with smoothed image to attract curve
 - ▶ Finish with unsmoothed image to get details

Scale

- In the simplest snakes algorithm, image features only attract locally
- Greater region of attraction: smooth image
 - ▶ Curve might not follow high-frequency detail
- Multi-resolution processing
- Looking for global minimum vs. local minima
 - ▶ Start with smoothed image to attract curve
 - ▶ Finish with unsmoothed image to get details

A Greedy Algorithm

- A greedy algorithm for deformable model
 - ▶ Makes locally optimal choices and hopes to lead to a globally optimal solution
 - ▶ Simplicity:
 - ★ Knowledge of the calculus of variations is not required
 - ▶ Low computational complexity:
 - ★ The number of iterations to converge proportional to the number of movement in each iteration times the number of contour points
- The energy functional is given by

$$\epsilon = \int (\alpha(s)E_{cont} + \beta(s)E_{curv} + \gamma(s)E_{image})ds$$


A Greedy Algorithm

- A greedy algorithm for deformable model
 - ▶ Makes locally optimal choices and hopes to lead to a globally optimal solution
 - ▶ Simplicity:
 - ★ Knowledge of the calculus of variations is not required
 - ▶ Low computational complexity:
 - ★ The number of iterations to converge proportional to the number of movement in each iteration times the number of contour points
- The energy functional is given by

$$\epsilon = \int (\alpha(s)E_{cont} + \beta(s)E_{curv} + \gamma(s)E_{image})ds$$

The Energy Functional


- The three energy terms:
 - ▶ Internal energy: E_{cont} and E_{curv} are for continuity and smoothness of the snake
 - ▶ External energy: E_{image} is for edge attraction
- α, β, γ control the relative influence of the energy term (can vary along the curve)



$$\epsilon = \int (\alpha(s)E_{cont} + \beta(s)E_{curv} + \gamma(s)E_{image})ds$$

The Energy Functional


- The three energy terms:
 - ▶ Internal energy: E_{cont} and E_{curv} are for continuity and smoothness of the snake
 - ▶ External energy: E_{image} is for edge attraction
- α, β, γ control the relative influence of the energy term (can vary along the curve)



$$\epsilon = \int (\alpha(s)E_{cont} + \beta(s)E_{curv} + \gamma(s)E_{image})ds$$

The Energy Terms

- Continuity term: $E_{cont} = \|\mathbf{p}_i - \mathbf{p}_{i-1}\|^2$
 - ▶ To make equal space between points
- Smoothness term: $E_{curv} = \|\mathbf{p}_{i-1} - 2\mathbf{p}_i + \mathbf{p}_{i+1}\|^2$
 - ▶ To avoid oscillation, penalize high contour curvature
- Edge attraction term: $E_{image} = -\|\nabla I\|$
 - ▶ Computed at each snake point, becomes very small (negative) near image edges (large gradient)



Snake Problem

- Given N initial points $\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N$, representing the initial position of the snake, fit the target image contour by minimizing the energy functional

$$\sum_{i=1}^N (\alpha_i E_{cont} + \beta_i E_{curv} + \gamma_i E_{image})$$

- Greedy minimization
 - The neighborhood over which the energy functional is locally minimized is typically small (3×3 or 5×5 window)
 - The local minimization is done by direct comparison of the energy functional values at each location
- Corner elimination
 - If a curvature maximum is found at point \mathbf{p}_i , then set β_i to zero to make the deformable contour piecewise smooth

Snake Problem

- Given N initial points $\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N$, representing the initial position of the snake, fit the target image contour by minimizing the energy functional

$$\sum_{i=1}^N (\alpha_i E_{cont} + \beta_i E_{curv} + \gamma_i E_{image})$$

- Greedy minimization
 - The neighborhood over which the energy functional is locally minimized is typically small (3×3 or 5×5 window)
 - The local minimization is done by direct comparison of the energy functional values at each location
- Corner elimination
 - If a curvature maximum is found at point \mathbf{p}_i , then set β_i to zero to make the deformable contour piecewise smooth

Snake Problem

- Given N initial points $\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N$, representing the initial position of the snake, fit the target image contour by minimizing the energy functional

$$\sum_{i=1}^N (\alpha_i E_{cont} + \beta_i E_{curv} + \gamma_i E_{image})$$

- Greedy minimization
 - The neighborhood over which the energy functional is locally minimized is typically small (3×3 or 5×5 window)
 - The local minimization is done by direct comparison of the energy functional values at each location
- Corner elimination
 - If a curvature maximum is found at point \mathbf{p}_i , then set β_i to zero to make the deformable contour piecewise smooth

Remarks


- The values of the parameters α_i , β_i and γ_i can be all set to 1, or $\alpha_i = \beta_i = 1, \gamma_i = 1.2$ (more weight on edge attraction)
- To prevent noisy corner:
 - ▶ A point is a corner if and only if the curvature is locally maximum at that point, and
 - ▶ The norm of the intensity gradient at that point is sufficiently large

Diffusion-Based Methods

- Another way to attract curve to localized features: vector flow or diffusion methods
- Example:
 - ▶ Find edges using Canny
 - ▶ For each point, compute distance to nearest edge
 - ▶ Push curve along gradient of distance field

Gradient Vector Fields

Check <http://iac1.ece.jhu.edu/projects/gvf/> for more details.



Gradient Vector Fields

Simple Snake

With Gradient Vector Field

Gradient Vector Fields